

# Concurrent programming (part 3)

# Agenda

- Atomic operations
- Concurrency utilities
- Testing concurrent applications

# Atomic operations

# Compare-and-Swap (CAS)

- Compare-and-swap is an atomic instruction that compares the location in memory with a given value and only if they are equal sets a new value
- CAS is used to implement atomic operations that achieve **optimistic locking** and are thread-safe

```
// represented by the following pseudocode
if (memoryLocation != value) {
    return false;
}
memoryLocation = newValue;
return true;
```

# Compare-and-Swap (CAS)

- The JDK provides support for calling CAS instructions through native code as provided by the **`jdk.internal.misc.Unsafe`** class
- An alternative of Unsafe as of JDK 9 that provides support for CAS is provided by the **`java.lang.invoke.VarHandle`** class

```
Unsafe.compareAndSetInt (...)
```

```
VarHandle.compareAndSet (...)
```

**Apart from performing CAS operations the Unsafe class also provides the possibility to access off-heap memory but since it is an internal class it is encapsulated as of JDK 9**

# Atomic variables

- Maintaining a single variable that is updatable from many threads is a common scalability issue
- Atomic variables already present in the JDK serve as a means to implement updatable variables in a multithreaded environment
- Atomic variables are part of the **java.util.concurrent.atomic** package

# Atomic variables

- They make use of the CAS support provided by the JDK to provide performant thread-safe operations over shared variables
- The other utilities that we already discussed that make use of CAS are concurrent collections like **ConcurrentHashMap**

# Atomic variables

- Atomic variables are provided by the following classes:
  - AtomicBoolean
  - AtomicInteger
  - AtomicIntegerArray
  - AtomicLong
  - AtomicLongArray
  - AtomicReference
  - AtomicReferenceArray
  - DoubleAccumulator
  - DoubleAdder
  - LongAccumulator
  - LongAdder



# Atomic variables

```
AtomicInteger value = new AtomicInteger();  
Thread thread = new Thread( () -> {  
    value.getAndIncrement();  
} );  
thread.start();  
value.getAndAdd(10);  
thread.join();  
System.out.println(value.get());
```

# Atomic variables

```
DoubleAccumulator accumulator =  
    new DoubleAccumulator((x, y) -> x + y , 0);  
Thread thread = new Thread( () -> {  
    accumulator.accumulate(0.9);  
} );  
thread.start();  
accumulator.accumulate(10.1);  
thread.join();  
System.out.println(accumulator.get());
```

# Atomic variables

```
AtomicReference<String> reference =  
    new AtomicReference<String>("first");  
Thread thread = new Thread( () -> {  
    reference.getAndAccumulate(" someText",  
        (x, y) -> {return x + y;});  
    } );  
thread.start();  
reference.getAndAccumulate(" otherText",  
    (x, y) -> {return x + y;});  
thread.start();  
thread.join();  
System.out.println(reference.get());
```

# Concurrency utilities

# Fork/Join framework

- **ForkJoinPool** is an implementation of the **ExecutorService** interface introduced in JDK 7
- Provides the possibility to execute tasks that can be organized in a divide-and-conquer manner
- Tasks submitted to the **ForkJoinPool** are represented by a **ForkJoinTask** instance

# Fork/Join framework

- Typical implementations of parallel tasks do not extend directly **ForkJoinTask**
- **RecursiveTask** instances can be used to execute parallel tasks that return a result
- **RecursiveAction** instances can be used to execute parallel tasks that do not return a result

# Fork/Join framework

- A global **ForkJoinPool** instance is used for any ForkJoinTasks that are not submitted to a particular **ForkJoinPool**
- To get a reference to the global ForkJoinPool instance the static **ForkJoinPool.commonPool()** method can be used
- This is the case with **parallel streams**: they make use of the common ForkJoinPool for task execution

# Fork/Join framework

- A global **ForkJoinPool** instance is used for any ForkJoinTasks that are not submitted to a particular **ForkJoinPool**
- To get a reference to the global ForkJoinPool instance the static **ForkJoinPool.commonPool()** method can be used
- This is the case with **parallel streams**: they make use of the common ForkJoinPool for task execution



# Fork/Join framework

```
public class NumberFormatAction extends RecursiveAction {
    ...
    @Override
    protected void compute() {
        if (end - start <= 2) {
            System.out.println(start + " " + end);
        } else {
            int mid = start + (end - start) / 2;
            NumberFormatAction left =
                new NumberFormatAction(start, mid);
            NumberFormatAction right =
                new NumberFormatAction(mid + 1, end);
            invokeAll(left, right);
        }
    }
}
```

# Fork/Join framework

```
public static void main(String[] args) {  
    NumberFormatAction action =  
        new NumberFormatAction(1, 50);  
    ForkJoinPool.commonPool().invoke(action);  
}
```

# Parallel streams

- Parallel streams can be created as regular streams with the **parallelStream** method

```
list.parallelStream()
```

- Parallel streams should be used only for time-intensive tasks rather than IO
- In many cases regular streams outperform parallel ones so they need to be used with caution

**When using parallel streams always measure performance with standard streams**

# CompletableFuture

- Provides a facility to create a chain of dependent non-blocking tasks
- An asynchronous task can be triggered as the result of a completion of another task



- A CompletableFuture may be completed/cancelled by a thread prematurely

# CompletableFuture

- Provides a very flexible API that allows additionally to:
  - combine the result of multiple tasks in a CompletableFuture
  - provide synchronous/asynchronous callbacks upon completion of a task
  - provide a CompletableFuture that executes when first task in group completes
  - provide a CompletableFuture that executes when all tasks in a group complete

# CompletableFuture

```
CompletableFuture<Integer> task1 = new  
    CompletableFuture<Integer>();  
  
// forcing completing of future by  
specifying result  
task1.complete(10);
```

# CompletableFuture

```
CompletableFuture<Integer> task1 =  
    CompletableFuture.supplyAsync(  
        () -> { ... return 10; });  
  
// executed on completion of the future  
task1.thenApply((x) -> {...});  
  
// executed in case of exception or completion  
// of the future  
task1.handle((x, y) -> {...});  
  
// can be completed prematurely with a result  
// task1.complete(20);  
System.err.println(task1.get());
```

# CompletableFuture

```
CompletableFuture<Object> prev = null;
Supplier<Object> supplier = () -> { ... };

for (int i = 0; i < count; i++) {
    CompletableFuture<Object> task;
    if (prev != null) {
        task = prev.thenCompose(
            (x) -> { return CompletableFuture
                .supplyAsync(supplier); });
    } else {
        task = CompletableFuture
            .supplyAsync(supplier);
    }
    prev = task;
}
prev.get();
```



# A few more things ...

- There is a **Timer** utility that can be used for scheduling tasks but a scheduled executor thread pool is more preferable as a more robust alternative
- A **ThreadLocalRandom** utility is provided since JDK 7 that can be used to generate random numbers for the current thread
- A **Flow** class introduced in JDK 9 provides a reactive streams specification for the JDK that can be used to provide a publish-subscribe

# Testing concurrent applications

# Testing concurrent applications

- Testing concurrent applications for correctness is inherently difficult
- In many cases non-deterministic unit tests are written that try to simulate running a particular piece of code in a multithreaded manner
- Testing frameworks provide certain utilities that can be used to facilitate testing of concurrent applications

# Testing concurrent applications

- A classic way of running multiple threads against code-under-test in a unit test is to use `CountDownLatch`:

```
@RepeatedTest(10)
public void testLinkedList()
    throws InterruptedException {
    CountDownLatch latch = new CountDownLatch(100);
    for(int i = 0; i < 100; i++) {
        new Thread( () -> {
            // unit under test ...
            latch.countDown();
        }).start();
    }
    latch.await();
    // perform asserts
}
```

# Testing concurrent applications

- Frameworks like **ThreadWeaver** provide the possibility to interleave execution threads
- Interleaving happens using breakpoints (line by line) so that the unit-under-test is tested using different thread ordering
- For performance testing a framework like **JMH** can be used to performs proper application warmup before the test is executed

Questions ?